Document Extraction Test Fixture

This PDF is synthetic. It exists so the PDF extractor has a real multi-page document with a genuine text layer to parse.

Nothing here describes a real person, organisation or event, and no third-party content is embedded. Only the base-14 Helvetica font is referenced, so no font data is redistributed.

The extractor falls back to Vision OCR when a document yields fewer than fifty characters, so this page deliberately carries several paragraphs of ordinary prose.

Page Two

A second page exists so that page counting is exercised rather than assumed. The tests assert a positive page count and a word count above ten.

Truncation is also checked: extracting with a small maximum size must return shortened text and set the truncation flag, which requires the full document to be comfortably longer than that limit.

Unicode is intentionally absent from this page so that the text operators stay within WinAnsi.

Page Three

The previous fixture in this slot was a real research report exported from a word processor. It carried author metadata and cited third-party writers by name.

This replacement was written for the repository and carries no third-party rights.